

Census Analysis

A Tech Talent South Production

Powered by Conor and Owen



Table of Contents



1. Overview

What? Who? Why?



Hello!


We are Conor and Owen!

Aspiring Data Scientists hoping
to bring some unique insights
using historical Census Data



Objective

- Provide high-level analysis of the Census data set
- Predict if an individual's income is more or less than \$50k/year



Data Overview

- Extracted from 1994 census bureau database
- 15 original data columns such as age, education, marital.status etc.
 - Transformed based on preliminary analysis (see Data Cleaning and Data Engineering slides)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 32561 non-null  int64
1   workclass           32561 non-null  object
2   fnlwgt              32561 non-null  int64
3   education           32561 non-null  object
4   education.num       32561 non-null  int64
5   marital.status      32561 non-null  object
6   occupation          32561 non-null  object
7   relationship        32561 non-null  object
8   race                32561 non-null  object
9   sex                 32561 non-null  object
10  capital.gain        32561 non-null  int64
11  capital.loss        32561 non-null  int64
12  hours.per.week      32561 non-null  int64
13  native.country      32561 non-null  object
14  income              32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country	income
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	4356	40	United-States	<=50K
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	4356	18	United-States	<=50K
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	4356	40	United-States	<=50K
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	0	3900	40	United-States	<=50K
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	0	3900	40	United-States	<=50K

2. Preliminary Analysis

Data Cleaning and Visualization



Data Cleaning

- “?” in ‘workclass’, ‘occupation’, and ‘native.country’ columns
 - Change to null values, then fill them with the mode (most commonly appearing value) for that column

```
data[data == '?'] = np.nan
for column in ['workclass', 'occupation', 'native.country']:
    data[column].fillna(data[column].mode()[0], inplace=True)
```

- Income is a categorical variable
 - ‘<=50K’ or ‘>50K’
- As the target variable, it needs to be numerical

```
#Replace the categorical variables with numerical variables
data['income'] = data['income'].replace({'<=50K':0, '>50K':1})
```


Data Visualization


GDS Overview




<https://datastudio.google.com/reporting/3a9c195c-3a28-4a99-bda7-6acbfa4a068f>

GDS Zoom In: Demographics


Sex Breakdown




Employment Breakdown




Marital Status




Income Breakdown



Race Breakdown

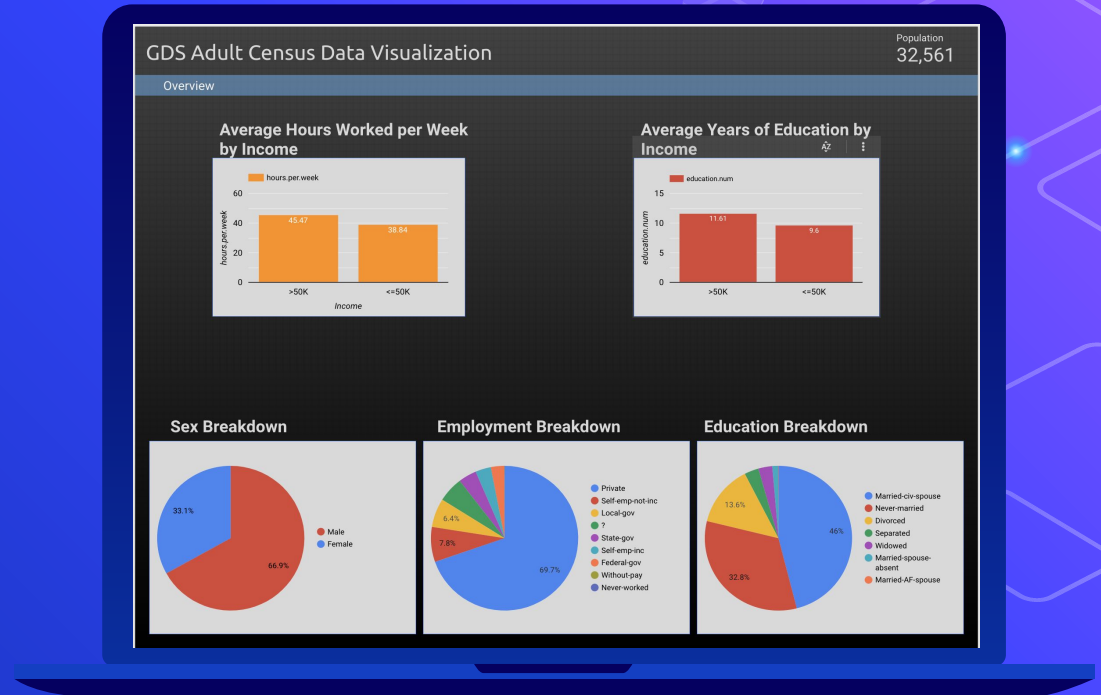


Education Breakdown



Data Visualization


GDS Overview 2



<https://datastudio.google.com/reporting/3a9c195c-3a28-4a99-bda7-6acbf4a068f>

GDS Zoom In: Income

Average Age by Income




Average Years of Education by Income



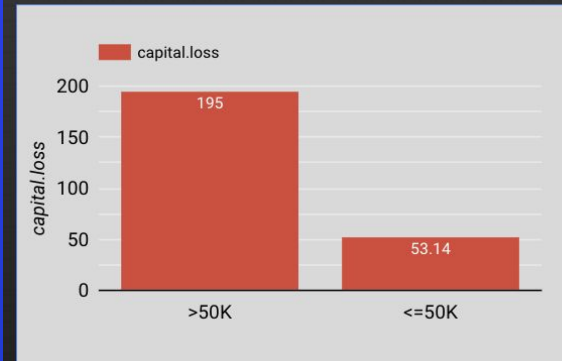
Average Hours Worked per Week by Income



Average Capital Gain by Income






Average Capital Loss by Income





Visuals - Categorical part 1

- Occupation vs Workclass
- Sex






Visuals - Categorical part 2

- Race
- Native Country



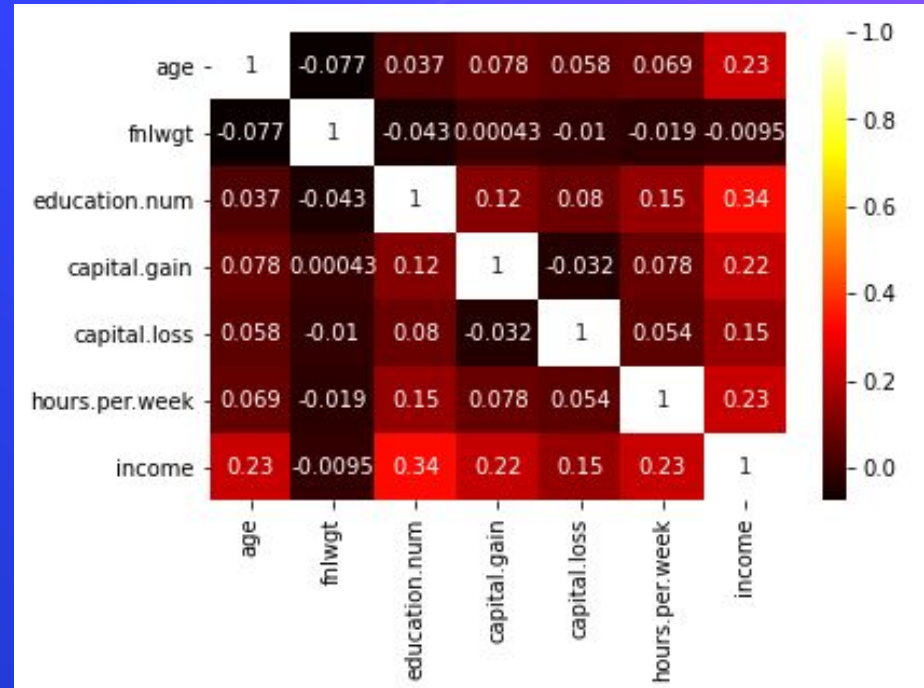
Visuals - Categorical part 2

- Relationship vs Marital Status






Visuals - Numerical part 1

- Education.num strong correlation (.34) with income
- Age, Capital.gain and Hours/week all decent indicators




Visuals - Numerical part 2

- Age, Hours/Week, Education Number



3. Data Engineering



Data Engineering

- Group 'relationship', 'race', 'sex', and 'occupation' by observed changes (from visualizations) and assign binary groups (0 or 1) to them
- Drop 'native.country', 'workclass', 'marital.status', 'education'
 - Redundant and uninformative columns

```
#Use .replace to set 0s and 1s
data['relationship'] = data['relationship'].replace(['Husband','Wife'], 1)\
                                                    .replace(['Not-in-family','Unmarried','Own-child','Other-relative'], 0)
data['race'] = data['race'].replace(['White','Asian-Pac-Islander'], 1)\
                    .replace(['Black','Other','Amer-Indian-Eskimo'], 0)
data['sex'] = data['sex'].replace('Male', 1)\
                    .replace('Female', 0)
data['occupation'] = data['occupation'].replace(['Exec-managerial', 'Prof-specialty','Adm-clerical', 'Sales',\
                                                'Tech-support'], 1)\
                    .replace(['Machine-op-inspct', 'Other-service', 'Craft-repair', 'Transport-moving',\
                              'Handlers-cleaners', 'Farming-fishing', 'Protective-serv', 'Armed-Forces',\
                              'Priv-house-serv'], 0)
data.drop(labels = ['native.country','workclass','marital.status','education'], axis = 1, inplace = True)
```

4.


ML Modeling

Data Cleaning and Visualization



“ All models are wrong,
but some are useful.

- George E. P. Box



Model Choices

NAIVE BAYES


- Requires predictors be independent
- Works well with binary classification data and many data points
- Fast to employ

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

N

LOGISTIC REGRESSION

- The dependent variable is binary, multinomial, or ordinal (most often binary)
- No multicollinearity in the model (independence tenet)



L


- Can handle both numerical and categorical data.
- High Variance in the data can create totally unique “trees” (results)

DECISION TREE

D



K



- Used for both classification and regression problems
- Groups by “neighbors”

K-NEAREST NEIGHBOR

Applying Models

- Using Naive Bayes, Logistic Regression, K-Nearest Neighbors and Decision Trees
- Test_size set to 25%
- Seed set to 0

```
#Set features
X = data[['occupation', 'relationship', 'race', 'sex', 'age', 'fnlwgt', 'capital.gain', 'capital.loss', 'hours.per.week']]
#Set target Variable
y = data['income']

#Set test/train split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)

#select the model
nb = GaussianNB()
#fit the model
nb.fit(X_train,y_train)
#predict based on the model
y_pred=nb.predict(X_test)
```

How to measure a model

⬡ TP = True Positive, TN = True Negative, FP = False Positive, FN = False Negative

⬡ **Accuracy**

- $(TP + TN) / (TP + TN + FP + FN)$
- Compares accurate predictions vs total # of predictions
- Describes overall accuracy of the model
- Sample = Entire dataset


⬡ **Precision**

- $(TP) / (TP + FP)$
- Compares accurate positive predictions vs all positive predictions
- Describes a model's accuracy when only considering positive predictions made
- Sample = All positives predicted

⬡ **Recall**

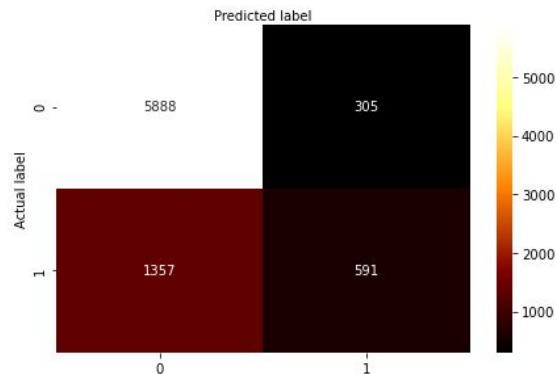
- $(TP) / (TP + FN)$
- Compares accurate positive predictions vs the true total of positive values
- Describes model's accuracy to predict positives in relation to the entire dataset
- Sample = All positives in the dataset

5. Results




Accuracy: 0.7958481758997666
Precision: 0.6595982142857143
Recall: 0.303388090349076

Naive Bayes Confusion matrix




Accuracy: 0.8003930721041641
Precision: 0.5800694100148736
Recall: 0.6006160164271047

Decision Tree Confusion matrix




Accuracy: 0.7991647217786513
Precision: 0.7158620689655173
Recall: 0.26642710472279263

Logistic Regression Confusion matrix



Accuracy: 0.7780370961798305
Precision: 0.5631154879140555
Recall: 0.3228952772073922

K-Nearest Neighbors Confusion matrix



Comparing Models

- Using our handmade, binary numeric variables

Model	Accuracy	Precision	Recall
Naive Bayes	79.58%	65.95%	30.33%
Decision Tree	80.03%	58.00%	60.06%
Logistic Regression	79.91%	71.58%	26.64%
KNN	77.80%	56.31%	32.28%


Applying sklearn preprocessing

- Easy way of turning categorical data into numerical, and standardizing it.

```
#create a list of our categorical columns for our for loop to iterate over
cats = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
#set our sklearn LabelEncoder to a variable
label_encoder = LabelEncoder()
#For each column from our list, fit the LabelEncoder and then transform the column as such
for column in cats:
    label_encoder.fit(data2[column])
    data2[column] = label_encoder.transform(data2[column])
#set our sklearn StandardScaler to a variable (this is like standardizing with z-scores, applied to all our columns)
scaler = StandardScaler()
#set our train/test split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state = 0)
#apply the scaler to the columns in test and train
x_train = pd.DataFrame(scaler.fit_transform(x_train), columns = x.columns)
x_test = pd.DataFrame(scaler.transform(x_test), columns = x.columns)
```


Accuracy: 0.7990418867461
Precision: 0.6788990825688074
Recall: 0.30390143737166325

Naive Bayes Label Encoded Confusion matrix




Accuracy: 0.8109568849035745
Precision: 0.6041772796739684
Recall: 0.608829568788501

Decision Tree Label Encoded Confusion matrix




Accuracy: 0.8027269377226385
Precision: 0.6993006993006993
Recall: 0.3080082135523614

Logistic Regression Label Encoded Confusion matrix



Accuracy: 0.823731728288908
Precision: 0.6480092325447201
Recall: 0.5764887063655031

K-Nearest Neighbors Label Encoded Confusion matrix



Label Encoding

- Using sklearn's LabelEncoder and StandardScaler functions for categorical variable transformation.
- Applying the same models to this new dataset

Model	Accuracy	Precision	Recall
Naive Bayes	79.90% ↑	67.88% ↑	30.39% ↑
Decision Tree	81.09% ↑↑	60.41% ↑	60.88% ↑
Logistic Regression	80.27% ↑	69.93% ↓	30.80% ↑
KNN	82.37% ↑↑↑	64.80% ↑↑↑	57.64% ↑↑↑

No single statistic
tells the whole story.



Changing the Seed

- Seed originally set to 0
 - Iterating through seeds 0-9 for each model

```
#Set test/train split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)
```


Naive Bayes



Naive Bayes

Logistic Regression

A dark blue laptop is centered in the frame. Its screen is black and displays the text "Logistic Regression" in a white, sans-serif font. The laptop is set against a blue background with faint, white, geometric line patterns and small blue dots. To the right of the laptop, there are three small, light blue speech bubble icons containing the binary strings "001", "011", and "010" from top to bottom.

Logistic Regression

Decision Tree

A stylized illustration of a laptop with a dark blue frame. The screen is black and displays the text "Decision Tree" in a white, sans-serif font. The background of the entire slide is a vibrant blue with faint, white, geometric line patterns and small, glowing blue dots. Some of these patterns resemble binary code or network connections, with small white boxes containing binary digits like "001", "011", and "010" scattered throughout.

Decision Tree

<https://media.giphy.com/media/ZVpzkB8yWMxjm33SOn/giphy.gif>

K-Nearest Neighbors

A stylized illustration of a laptop with a dark blue frame. The screen is black and displays the text 'K-Nearest Neighbors' in white. The background of the entire slide is a vibrant blue with a network of white lines and nodes, some of which are labeled with binary code (001, 011, 010).



K-Nearest Neighbors

Results




- Most Accurate - KNN
- Most Precise - Logistic Regression
- Highest Recall - Decision Tree

Summary

- Since there is neither a high cost associated with False Negatives nor False Positives, the best model to use to predict an individual's Income using the 1994 Census data is KNN.
- We are able to predict with over **80%** accuracy whether or not someone will make **greater than or less than \$50k**
- Capital Gain, Age, and Hours Worked per Week were the **strongest predictors of Income**. All were **positively** correlated with **>\$50k Income** (i.e. as capital.gain/age/hours.worked increased so did likelihood of earning greater than \$50k)



Appendix



Credits

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by [SlidesCarnival](https://www.slidescarnival.com/aliena-free-presentation-template/4597#preview) designed by Jimena Catalina
- Photographs by [Unsplash](https://unsplash.com/)
- This data was extracted from the 1994 Census bureau database by Ronny Kohavi and Barry Becker (Data Mining and Visualization, Silicon Graphics)
- Ron Kohavi, "Scaling Up the Accuracy of Naive-Bayes Classifiers: a Decision-Tree Hybrid", Proceedings of the Second International Conference on Knowledge Discovery and Data Mining, 1996. (PDF)
- Kaggle inspiration: <https://www.kaggle.com/uciml/adult-census-income>
- George Box image: [https://en.wikipedia.org/wiki/File:GeorgeEPBox_\(cropped\).jpg](https://en.wikipedia.org/wiki/File:GeorgeEPBox_(cropped).jpg)




SlidesCarnival icons are editable shapes.

This means that you can:

- Resize them without losing quality.
- Change fill color and opacity.
- Change line color, width and style.

Isn't that nice? :)

Examples:



Find more icons at
slidescarnival.com/extra-free-resources-icons-and-maps

Thank you!

Any questions?


You can find us at:

owinters58@gmail.com

conoranderson2@gmail.com



Diagrams and infographics



You can also use any emoji as an icon!
And of course it resizes without losing quality.

How? Follow Google instructions

<https://twitter.com/googledocs/status/730087240156643328>



many more...